

Leveraging Large Language Models (LLM) for AIS Vessel Trajectory Prediction

CSCI 6917 - Guided Research Grad I

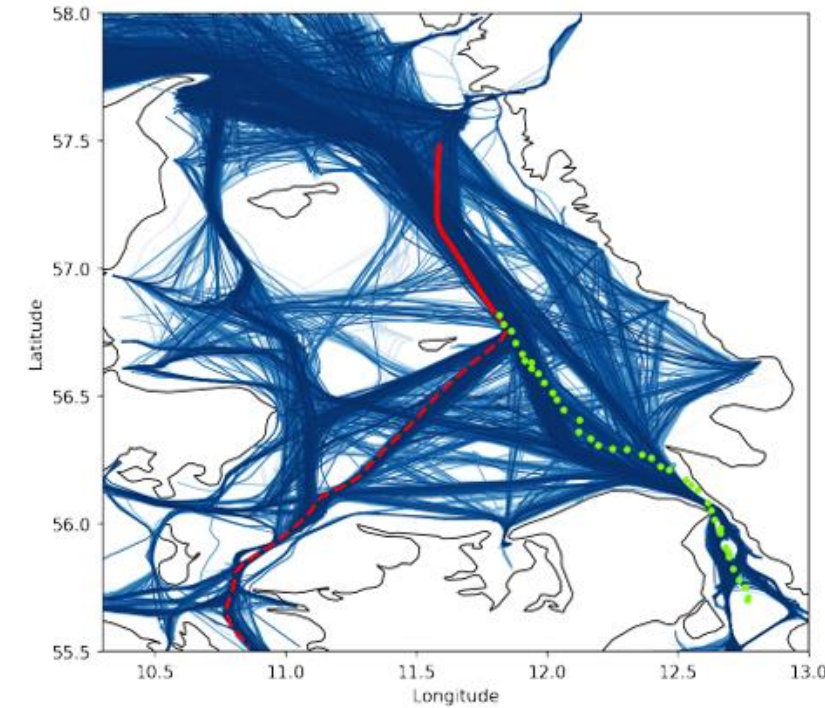
Ryan Gross (G47667332)

Summer Semester - 2023

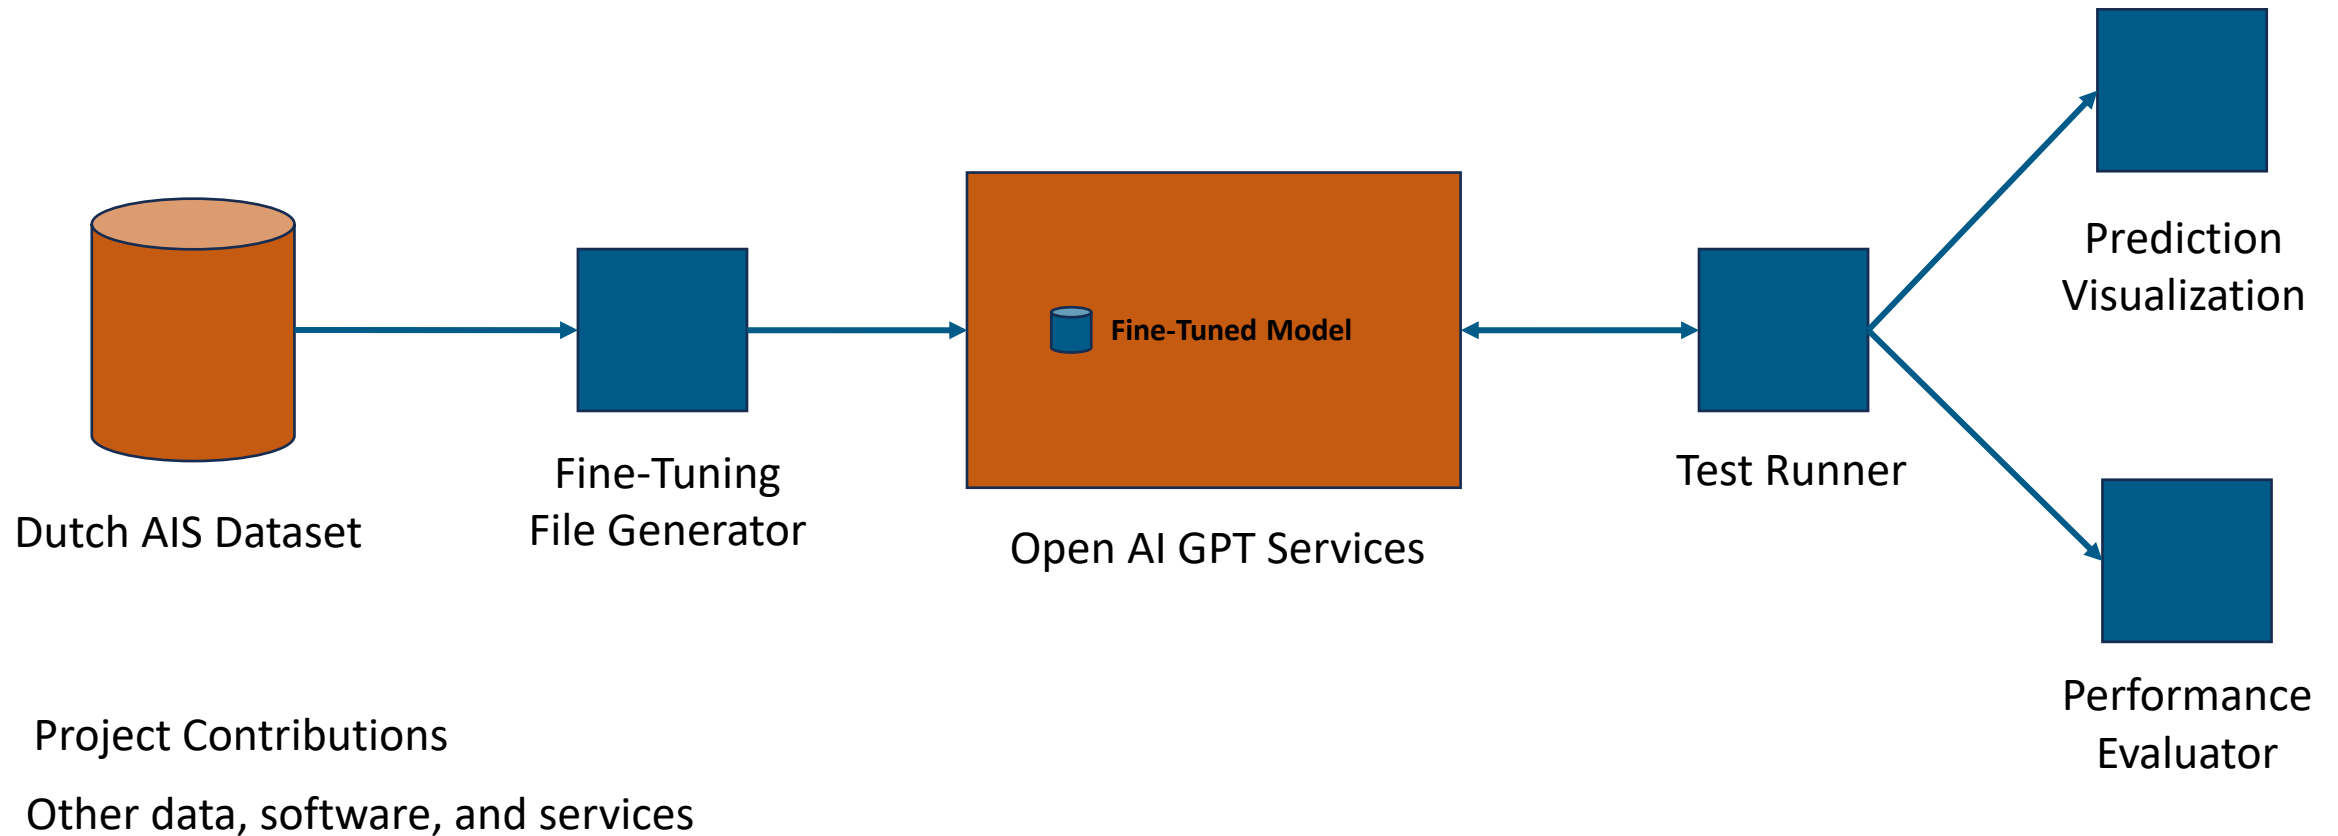


Project Objective

- What are you trying to do? Articulate your objectives using absolutely no jargon.
 - **Predict the future trajectory of cargo vessels.**
- How is it done today, and what are the limits of current practice?
 - **Various seq2seq approaches such as RNN, LSTM, and transformers are used, but implementation is complex.**
- What is new in your approach and why do you think it will be successful?
 - **The approach treats the problem as a language translation task, which LLMs are good at, allowing pre-trained LLMs to solve the problem with little effort and complexity from the software engineer.**
- Who cares? If you are successful, what difference will it make?
 - **EMS, Coast Guard, Navy, and other maritime forces could better locate distressed or malign vessels.**
- What are the risks?
 - **LLM may not properly format messages, LMM may not pick up on vessel patterns.**
- How much will it cost? How long will it take?
 - **\$12 and a summer semester.**
- What are the mid-term and final “exams” to check for success?
 - **Mid-term: Ensure LLM can issue predictions with the correct data format**
 - **Final: Performance beats or is near state-of-the-art**




Technical Approach - Architecture



Technical Approach – Key Steps

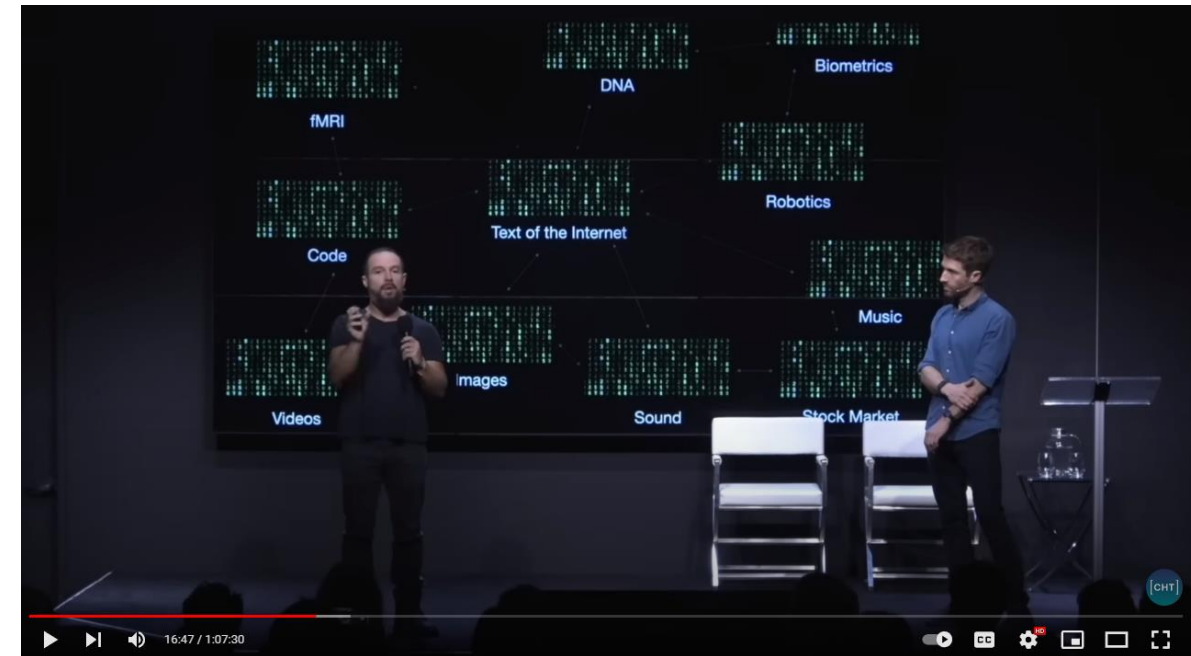
1. Find and gather dataset – Needed to be free and large.
2. Clean data – Remove paths over land, non-moving vessels, impossibly fast trips, etc
3. Find best base model – Run Few-Shot learning on each model to determine best model to fine-tune
 1. Format data using prompt engineering
 2. Evaluate
 3. Visualize
 4. Compare
4. Create trained model – Create custom model by fine-tuning best base model
 1. Format data in GPT required JSONL for training
 2. Ensure time and cost of training meets requirements/limits
 3. Run fine-tuning job
5. Evaluate trained model – Run test set and measure accuracy loss
6. Visualize – Plotting
 1. Individual predictions to better understand results
 2. Overall accuracy loss across test set



Technical Approach – Innovation

- This research further investigates the intuition that a large portion of prediction problems can be thought of as seq2seq, generative language problems.
- Examples:
 - Text phrase -> Python code (Code generation)
 - Text phrase -> Image (Image generation)
 - Image -> Next image (Image sequence generation)
 - English phrase -> French phrase (Speech translation)
 - Research paper text -> Research paper abstract (Text summarization)
 - Stock market historical data -> Stock market buy & sell tasks
 - Python code -> Code comments and documentation
 - Historical GPS locations -> Future GPS locations (This project)
- Leveraging existing, expertly trained models, could reduce the cost and complexity of solving many prediction problems.

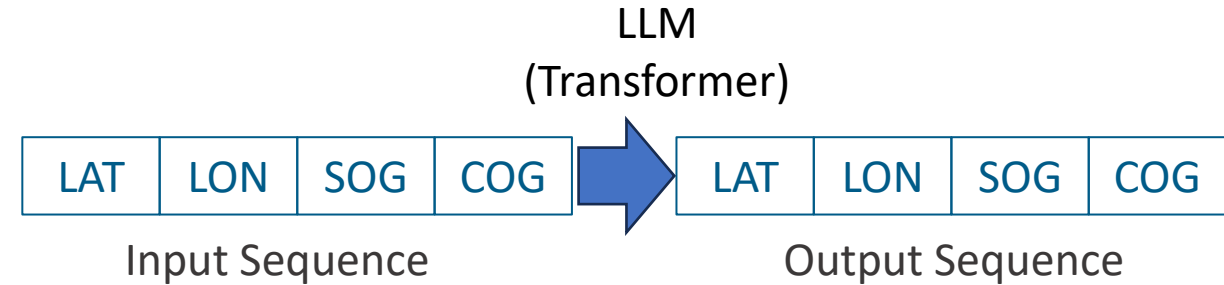
“Treat everything as language”



<https://youtu.be/xoVJKj8lcNQ?t=967>

Technical Approach – Innovation Cont.

- The dataset is prepared to contain four input fields:
 1. Latitude
 2. Longitude
 3. Speed over Ground (SOG)
 4. Course over Ground (COG)
- The first half of a vessel trip is the input to the model, and the second half of the trip is the output the model should predict.
- After enough input output pairs, the LLM learns to complete the vessel's trip only given the input. This is similar to text completion in email or texts.
- Why LLMs instead of just using a transformer?
 - Training a transformer from scratch requires learning both data format and data content patterns.
 - LLMs are already trained to detect text format patterns. This is why “fine-tuning” LLMs has become popular.




```
INPUT:
[[[0.902284800000001, 0.000241481481481024, 0.33666666666666667, 0.1908333333333335],
[0.9062707999999958, 0.01855555555555442, 0.35, 0.2461111111111109],
[0.9075819201171441, 0.03881266653955673, 0.33699999999999997, 0.2066666666666667],
[0.9119572001057434, 0.05661437692847809, 0.346969696969697, 0.189368686868685],
[0.9168839999999989, 0.07477777777777721, 0.35, 0.17],
[0.9206420000000008, 0.09309074074074032, 0.3433333333333334, 0.2525]]]


OUTPUT:
[[[0.9207287999999977, 0.11288074074074035, 0.32, 0.24000000000000002],
[0.9204009092225505, 0.1327053198908483, 0.3560606060606061, 0.2618181818181818],
[0.9197055001310531, 0.15264774695520808, 0.33999999999999997, 0.26296296296296295],
[0.9175780667928108, 0.1717417593857311, 0.3430555555555556, 0.2952083333333335],
[0.9138304001005337, 0.1907566976679734, 0.3461111111111111, 0.30344907407407407],
[0.909522254634959, 0.208258417651993, 0.31969696969696965, 0.31782828282828285]]]
```

Results


GPT 3.0
Curie
Few Shot



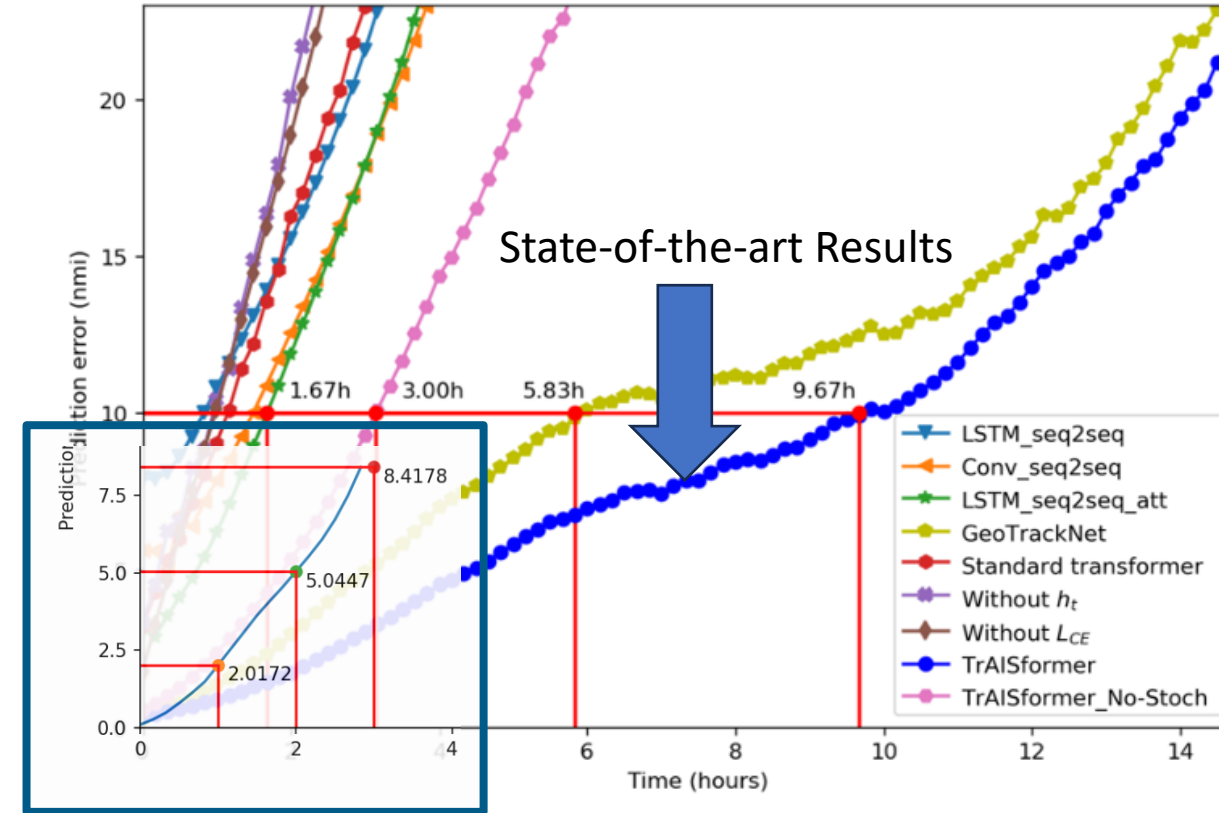

GPT 4.0
Few Shot



GPT 3.0
Davinci
Few Shot



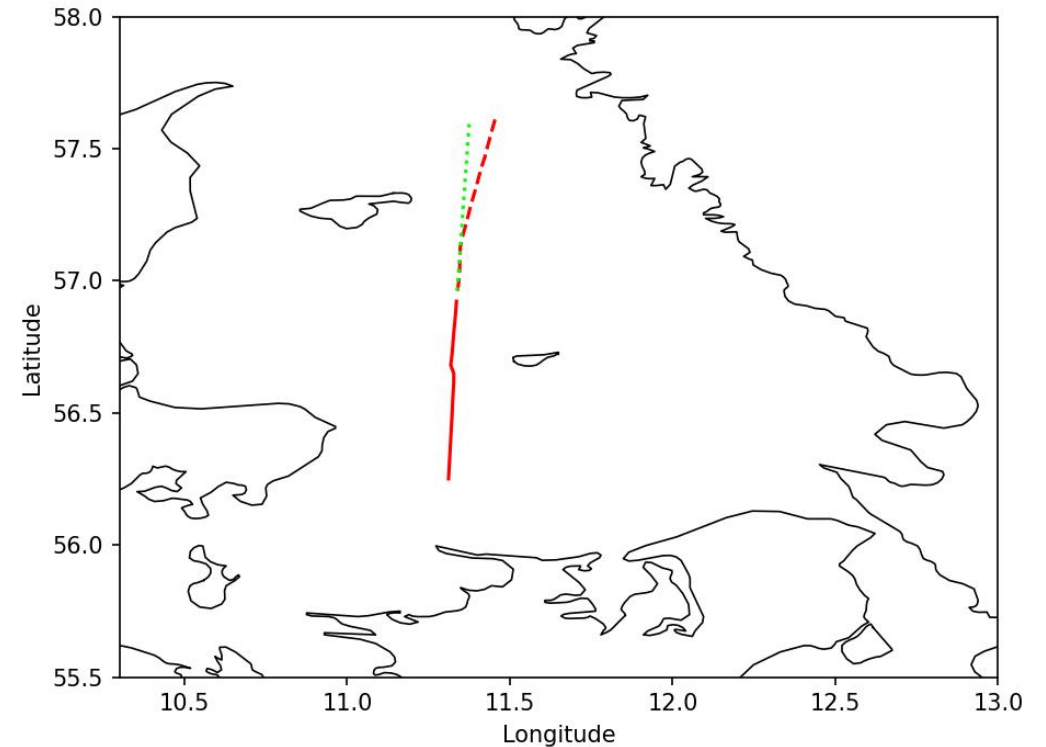
GPT 3.0
Davinci
Fine-Tuned



Best Results from Project
(GPT 3.5 Turbo, Few Shot)

Conclusion

- GPT 3.5 Turbo
 - With very few lines of code and “few-shot” training, GPT 3.5 Turbo outperforms many non-transformer algorithms.
 - GPT 3.5 Turbo seemed to perform about as well as a human, drawing a strait line from the end of the input.
 - GPT 3.5 Turbo picked up on LAT and LON patterns but did not pick up on SOG & COG patterns.
 - Fine-tuning not available
- GPT 3.0
 - Accepts too few characters as input and output to effectively pick up on patterns.
 - Training is expensive. Original 26.3MB training file had to be reduced to 177KB to stay under \$10 billing limit.
- GPT 4.0
 - Strangely, performed about as well as GPT 3.0.
 - Fine-tuning not available
- While the original intuition and approach remains valid, OpenAI’s services allow for very limited customization or fine-tuning. Additionally, the cost to fine-tune is prohibitive to research students.



GPT 3.5 Turbo few-shot prediction in green, actual is dashed red, input is solid red.

Future Work

- Graduate Thesis Level
 - Continue to work with GPT 4.0 prompt engineering to try to increase performance.
 - Train an array of open source LLMs to compare results to state-of-the-art.
 - Compare LLMs with transformers trained from scratch.
 - Try different combinations of input and output features.
- PhD Dissertation Level
 - Understand how the transformer is learning trajectory patterns.
 - Understand how to optimize transformer/LLM parameters for trajectory prediction (heads, layers, embedded layers, learning rate, warm up tokens).
 - Build and demonstrate a transformer/LLM model that is state-of-the-art in location prediction.
 - Demonstrate applicability and versatility of model on ground, maritime, and air domains, given domain specific features.

References

- [1] Vaswani, Ashish et al. “Attention is All you Need.” NIPS (2017).
- [2] Nguyen, Duong and Ronan Fablet. “TrAISformer-A generative transformer for AIS trajectory prediction.” ArXiv abs/2109.03958 (2021): n. pag..
- [3] D. Nguyen, R. Vadaine, G. Hajduch, R. Garello, and R. Fablet, “A Multi-task Deep Learning Architecture for Maritime Surveillance using AIS Data Streams,” in 2018 IEEE International Conference on Data Science and Advanced Analytics (DSAA), Oct. 2018.
- [3] D. Nguyen, R. Vadaine, G. Hajduch, R. Garello, and R. Fablet, “GeoTrackNet-A Maritime Anomaly Detector using Probabilistic Neural Network Representation of AIS Tracks and A Contrario Detection,” IEEE Transactions on Intelligent Transportation Systems, Feb. 2021.
- [5] Ouyang, Long et al. “Training language models to follow instructions with human feedback.” ArXiv abs/2203.02155 (2022): n. Pag.
- [6] Brown, Tom B. et al. “Language Models are Few-Shot Learners.” ArXiv abs/2005.14165 (2020): n. Pag.